

1. Scatter Plots & Relationships among /proc/*/sched Statistics

I wrote the following awk script to create a scatter plot from any two variables in the in the /proc/*/sched files. `sched_scatter_plot.awk`:

```
BEGIN {
    num_processes = 0;

    if (x_axis=="") { x_axis = "sum_exec_runtime"; }
    if (y_axis=="") { y_axis = "vruntime"; }

    if (x=="") { x = 100; }
    if (y=="") { y = 30; }
}

FNR == 1 {
    num_processes++;
}

match($1, y_axis) { y_axis_vals[num_processes-1]=$3; next; }
match($1, x_axis) { x_axis_vals[num_processes-1]=$3; next; }

END {
    asort(y_axis_vals);
    asort(x_axis_vals);

    printf("Maximum %s: %f\n", y_axis, y_axis_vals[num_processes]);
    printf("Maximum %s: %f\n\n", x_axis, x_axis_vals[num_processes]);

    if (log_scale_x > 1 || log_scale_y > 1) {
        for (i=1; i<=num_processes; i++) {
            if (log_scale_x) {
                x_axis_vals[i] = log(x_axis_vals[i] + 1) / log(log_scale_x);
            } if (log_scale_y) {
                y_axis_vals[i] = log(y_axis_vals[i] + 1) / log(log_scale_y);
            }
        }
    }

    print y_axis;

    x_max = x_axis_vals[num_processes];
    y_max = y_axis_vals[num_processes];
    x_scale = (x_max > 0) ? x / x_max : 1;
    y_scale = (y_max > 0) ? y / y_max : 1;

    for (i=1; i<=num_processes; i++) {
        x_coord = int(x_scale * x_axis_vals[i]);
        y_coord = int(y_scale * y_axis_vals[i]);

        plot[x_coord, y_coord] += 1;
        if (plot[x_coord, y_coord] > 9) { plot[x_coord, y_coord] = 9; }
    }

    print " ^";
    for (curr_y=y; curr_y >= 0; curr_y--) {
        printf " |";
        for (curr_x=0; curr_x<=x; curr_x++) {
            printf (plot[curr_x, curr_y] ? plot[curr_x, curr_y] : " ");
        }
    }
}
```

```
    printf "\n";
}

printf "  +";
for (i=0; i<x; i++) { printf "-"; }
printf (> "%s\n", x_axis);
}
```

Note: I prompted ChatGPT with the prompt “Why isn’t the graph produced by this awk graph scaling correctly” in order to debug an error in my code.

This awk script was used to generate log-scaled graphs showing the relationship between the inputted variables. The numbers on the graphs represent how many processes had the corresponding relationship between the variables, i.e. how many dots there would be at any given position in a scatter plot.

Before generating the graphs, I restarted the virtual machine and waited approximately 10 minutes to make the data consistent with section 2.




Figure 1: ‘vruntime’ vs. ‘sum exec runtime’

Note: I now realize that I was supposed to make this graph showing $\log(vruntimes)$ vs sum_exec_runtime , not both log-scaled. However, it took forever to actually get the data from sections 1 and 2 to be consistent outside of the kernel parameters, so I am just going to leave a correct graph in the Appendix, sorry about that.

Figure 1 shows the relationship of `vruntime` and total CPU runtime of all the processes on my root machine. There is a relatively strong logarithmic correlation between `vruntime` and total CPU time. Figure 1 was produced using the following bash command:

```
awk -f sched_scatter_plot.awk -v x=100 -v y=40
-v x_axis=sum_exec_runtime
-v y_axis=vruntime
-v log_scale_x=2
-v log_scale_y=2
/proc/*sched
> vruntime_vs_sumexecetime.plot
```




Figure 2: 'voluntary_switches' vs. 'involuntary_switches'

Figure 2 shows the relationship of voluntary and involuntary context switches of all the processes on my root machine; there is a relatively strong direct correlation. It was produced using the following bash command:

Intro to OS Assignment 1

```
awk -f sched_scatter_plot.awk -v x=100 -v y=40
-v x_axis=nr_voluntary_switches
-v y_axis=nr_involuntary_switches
-v log_scale_x=2
-v log_scale_y=2
/proc/*sched
> voluntary_vs_involuntary_switches.plot
```



Figure 3: 'switches' vs. 'migrations'

Figure 3 shows the relationship of context switches and CPU migrations of all the processes on my root machine; there is a relatively strong direct relationship. It was produced using the following bash command:

```
awk -f sched_scatter_plot.awk -v x=100 -v y=40
-v x_axis=nr_switches
-v y_axis=se.nr_migrations
-v log_scale_x=2
-v log_scale_y=2
/proc/*sched
> switches_vs_migrations.plot
```

2. /proc/*/sched Statistics with Kernel Parameter Manipulation

For this section, I began by creating a bash script that increased all of the specified kernel parameters by a factor of 10:

```
sysctl -w kernel.sched_latency_ns=60000000
-w kernel.sched_migration_cost_ns=5000000
-w kernel.sched_min_granularity_ns=7500000
-w kernel.sched_nr_migrate=320
-w kernel.sched_rr_timeslice_ms=1000
-w kernel.sched_rt_period_us=10000000
-w kernel.sched_rt_runtime_us=9500000
-w kernel.sched_shares_window_ns=100000000
```

Note: the last two kernel parameters specified by the assignment (kernel.sched_shares_window_ns and kernel.sched_time_avg_ms) didn't exist on my root machine according to sysctl.

Then, I restarted the virtual machine and ran the bash script to change the parameters immediately. I then waited 10 minutes and generated the plots in the using the same script and commands as section 1:




Figure 4: 'vruntime' vs. 'sum_exec_runtime', multiplied kernel parameters

Intro to OS Assignment 1



Figure 5: ‘voluntary_switches’ vs. ‘involuntary_switches’ multiplied kernel parameters



Figure 6: ‘switches’ vs. ‘migrations’ multiplied kernel parameters

After increasing the kernel parameters tenfold, the relationship between voluntary and involuntary switches changed the most out of the three graphs created.

As figures 5 and 6 show, the number of context switches after multiplying the kernel parameters increased by a factor of about 10 from a maximum of 6532 to 61990 switches. Additionally, the number of involuntary_switches actually decreased despite the larger number of switches; this was likely due to the increased timeslice, which allowed processes to finish their work within a timeslice more often and give up their CPU time voluntarily more often.

Note that the spread of these graphs look very similar despite tenfold differences due to the log-scaling and differences in zooming/linear scaling.

Another notable comparison is that, although the number of switches increased dramatically, the number of migrations actually stayed somewhat constant, meaning there were 10x more switches per migration. This is probably due to increasing the cost of migration.

The relationship between vruntime and total CPU time stayed approximately the same after changing the kernel parameters.

3. Scatter Plots & Relationships for top Statistics

For this section, I began by enabling the CODE and DATA' fields in my.toprc configuration, then I used the following command to output the top data to a file for the awk script to read:

```
top -b -n 1 > top.out
```

Then, I manipulated my scatter_plot awk script from section 1 so that I could read data from the top output and specify the top fields that I want to read:

```
BEGIN {
    num_processes = 0;
    processing_lines_bool = 0;

    if (x_axis=="") { x_axis = "PID"; }
    if (y_axis=="") { y_axis = "TIME"; }

    if (x=="") { x = 100; }
    if (y=="") { y = 30; }
}

match($0, x_axis) && match($0, y_axis) {
    processing_lines_bool = 1;
    print $0;

    for (i=1; i<100; i++) {
        if (match($i, x_axis)) {
            x_col = i;
        }
        if (match($i, y_axis)) {
            y_col = i;
        }
    }

    next;
}

processing_lines_bool { num_processes++;
    x_axis_vals[num_processes]=$x_col;
    y_axis_vals[num_processes]=$y_col;
}

END {
    print num_processes;

    if (!processing_lines_bool) {
        print "Processes not found";
        exit;
    }

    asort(y_axis_vals);
    asort(x_axis_vals);

    print y_axis;
```

```
if (log_scale_x > 1 || log_scale_y > 1) {
    for (i=1; i<=num_processes; i++) {
        if (log_scale_x) {
            x_axis_vals[i] = log(x_axis_vals[i] + 1) / log(log_scale_x);
        } if (log_scale_y) {
            y_axis_vals[i] = log(y_axis_vals[i] + 1) / log(log_scale_y);
        }
    }
}

x_max = x_axis_vals[num_processes];
y_max = y_axis_vals[num_processes];
x_scale = (x_max > 0) ? x / x_max : 1;
y_scale = (y_max > 0) ? y / y_max : 1;

for (i=1; i<=num_processes; i++) {
    x_coord = int(x_scale * x_axis_vals[i]);
    y_coord = int(y_scale * y_axis_vals[i]);

    plot[x_coord, y_coord] += 1;
    if (plot[x_coord, y_coord] > 9) { plot[x_coord, y_coord] = 9; }
}

print " ^";
for (curr_y=y; curr_y >= 0; curr_y--) {
    printf " |";
    for (curr_x=0; curr_x<=x; curr_x++) {
        printf (plot[curr_x, curr_y] ? plot[curr_x, curr_y] : " ");
    }
    printf "\n";
}

printf " +";
for (i=0; i<x; i++) { printf "-"; }
printf (> %s\n", x_axis);

}
```




Figure 7: 'PID' vs. 'TIME'

Figure 10 shows the relationship between process ID and total CPU time of all the processes the machine. It was produced using the following bash command:

```
awk -f top_scatter_plot.awk -v x_axis="PID" -v y_axis="TIME" -v y=40  
-v log_scale_x=2 -v log_scale_y=2 top.out
```

The relationship between TIME and PID shows that (for the most part) only processes with very high process IDs last a very long time.




Figure 8: ‘CODE’ vs. ‘DATA’

Figure 11 shows the relationship between the size of executable's machine code and the corresponding amount of memory the process has allocated for all the processes on the machine. It was produced using the following bash command:

```
awk -f top_scatter_plot.awk -v x_axis="CODE" -v y_axis="DATA" -v y=40  
-v log scale x=2 -v log scale y=2 top.out
```

This relationship shows that the amount of machine code in a program and the amount of data that program uses is directly correlated.

4. CPU Migrations Animation

This was the C program that I wrote (i.e. my modification of c3.c) to create an animation showing which processes were on which CPU core.

Intro to OS Assignment 1

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main() {
    FILE *fp;
    char pid[10], user[32], pr[8], ni[8], virt[16], res[16], shr[16], s[4];
    char cpu[8], mem[8], time[32], command[128];
    unsigned p;

    char final_strings[4][2048];

    for (int i = 0; i < 10; i++) {
        memset(final_strings, 0, sizeof(final_strings));

        system("top -b -n 1 | tail -n +8 | sort > out.top");

        fp = fopen("out.top", "r");
        if (fp == NULL) {
            perror("Error opening file");
            return 1;
        }

        char buffer[1024];

        while (fgets(buffer, sizeof(buffer), fp) != NULL) {
            sscanf(buffer, " %s %d\n", pid, user, pr,
                   ni, virt, res, shr, s, cpu, mem, time, command, &p);
            char pid_buf[32];
            sprintf(pid_buf, "%s, ", pid);
            strncat(final_strings[p], pid_buf);
        }

        for (int p = 0; p < 4; p++) {
            printf("Core %d: %s\n\n", p, final_strings[p]);
        }

        fclose(fp);

        system("sleep 1");
        system("clear");
    }

    return 0;
}
```

The my changes from c3.c to this code are shown in the Appendix.

I compiled it and ran it to create the following output snapshots:

```

Core 0: 3, 4, 6, 8, 9, 11, 12, 14, 36, 103, 108, 112, 116, 120, 124, 128, 132, 136, 139, 145, 244, 262, 294, 296, 302, 303, 306, 310, 372, 401, 447, 495, 548, 550, 786, 798, 8  
26, 831, 961, 971, 1062, 1157, 1450, 1501, 1951, 1980, 2060, 2067, 2185, 2200, 2217, 98713, 102419, 480825, 1820353, 1845847, 1847694, 1849348, 1849376, 1850405, 1844343

Core 1: 10, 15, 16, 17, 18, 20, 34, 91, 94, 96, 109, 113, 117, 121, 125, 129, 133, 137, 140, 158, 231, 234, 238, 242, 246, 252, 254, 276, 277, 282, 284, 293, 307, 308, 309, 31  
1, 313, 315, 317, 348, 362, 535, 785, 789, 814, 819, 861, 938, 940, 948, 949, 950, 952, 959, 960, 1011, 1073, 1129, 1159, 1225, 1518, 1953, 2043, 2063, 2064, 2070, 2071, 2077,  
2093, 2102, 2108, 2111, 2114, 2135, 2183, 1842976, 1848599, 1849996, 1850029, 1850403,
```

Intro to OS Assignment 1

Core 0: 3, 4, 6, 8, 9, 10, 11, 12, 14, 36, 103, 108, 112, 116, 120, 124, 128, 132, 136, 139, 145, 244, 262, 294, 296, 302, 303, 306, 310, 372, 401, 447, 495, 548, 550, 786, 798, 826, 831, 961, 971, 1062, 1157, 1450, 1501, 1951, 1980, 2060, 2067, 2185, 2200, 2217, 98713, 102419, 480825, 1820353, 1845847, 1847694, 1849348, 1849376, 1850049, 1850472,

Core 1: 15, 16, 17, 18, 20, 34, 91, 94, 96, 109, 113, 117, 121, 125, 129, 133, 137, 140, 158, 231, 234, 238, 242, 246, 252, 254, 276, 277, 282, 284, 293, 307, 308, 309, 311, 313, 315, 317, 348, 362, 785, 789, 814, 819, 861, 938, 940, 948, 949, 950, 952, 959, 960, 1011, 1073, 1129, 1159, 1225, 1518, 1953, 2043, 2063, 2064, 2070, 2071, 2077, 2093, 2102, 2108, 2111, 2114, 2135, 2183, 1842970, 1848599, 1849966, 1850029, 1850470,

Core 2: 2, 21, 22, 23, 24, 26, 33, 35, 37, 38, 39, 92, 95, 97, 102, 106, 110, 114, 118, 122, 126, 130, 134, 138, 141, 143, 155, 220, 235, 240, 258, 260, 287, 288, 289, 292, 295, 297, 298, 299, 300, 304, 305, 353, 355, 787, 800, 817, 856, 876, 877, 927, 941, 958, 1093, 1163, 1193, 1198, 1332, 1497, 1505, 1510, 1511, 2176, 244914, 1011450, 1600039, 1841476, 1843936, 1850144, 1850157, 1850377, 1850473, 504,

Core 3: 1, 27, 28, 29, 30, 32, 40, 41, 42, 90, 93, 98, 105, 107, 111, 115, 119, 123, 127, 131, 135, 142, 171, 237, 239, 241, 243, 245, 247, 248, 249, 250, 251, 253, 255, 256, 257, 259, 261, 263, 264, 301, 312, 314, 316, 347, 349, 446, 577, 788, 808, 816, 857, 925, 1068, 1128, 1166, 1507, 1609, 1986, 1986, 1991, 2039, 2056, 2058, 2061, 2076, 2188, 102345, 55988, 1820354, 1841714, 1849300, 1849899, 1850014, 1850022, 1850023, 1850024, 1850145, 1850471,

Core 0: 3, 4, 6, 8, 9, 11, 12, 14, 36, 103, 108, 112, 116, 120, 124, 128, 132, 136, 139, 145, 244, 262, 294, 296, 302, 303, 306, 310, 372, 401, 447, 495, 548, 550, 786, 798, 826, 831, 961, 971, 1062, 1157, 1450, 1501, 1951, 1980, 2060, 2067, 2185, 2200, 2217, 98713, 102419, 480825, 1820353, 1845847, 1847694, 1849348, 1849376, 1850049, 1850472,

Core 1: 10, 15, 16, 17, 18, 20, 34, 91, 94, 96, 109, 113, 117, 121, 125, 129, 133, 137, 140, 158, 231, 234, 238, 242, 246, 252, 254, 276, 277, 282, 284, 293, 307, 308, 309, 311, 313, 315, 317, 348, 362, 785, 789, 814, 819, 861, 938, 940, 948, 949, 950, 952, 959, 960, 1011, 1073, 1129, 1159, 1225, 1518, 1953, 2043, 2063, 2064, 2070, 2071, 2077, 2093, 2102, 2108, 2111, 2114, 2135, 2183, 1842970, 1848599, 1849966, 1850029, 1850470,

Core 2: 2, 21, 22, 23, 24, 26, 33, 35, 37, 38, 39, 92, 95, 97, 102, 106, 110, 114, 118, 122, 126, 130, 134, 138, 141, 143, 155, 220, 235, 240, 258, 260, 287, 288, 289, 292, 295, 297, 298, 299, 300, 304, 305, 353, 355, 787, 800, 817, 856, 876, 877, 927, 941, 958, 1093, 1163, 1193, 1198, 1332, 1497, 1505, 1510, 1511, 2176, 244914, 1011450, 1600039, 1841476, 1843936, 1850144, 1850157, 1850377, 1850473, 504,

Core 3: 1, 27, 28, 29, 30, 32, 40, 41, 42, 90, 93, 98, 105, 107, 111, 115, 119, 123, 127, 131, 135, 142, 171, 237, 239, 241, 243, 245, 247, 248, 249, 250, 251, 253, 255, 256, 257, 259, 261, 263, 264, 301, 312, 314, 316, 347, 349, 446, 577, 788, 808, 816, 857, 925, 1068, 1128, 1166, 1507, 1609, 1986, 1986, 1991, 2039, 2056, 2058, 2061, 2076, 2188, 102345, 55988, 1820354, 1841714, 1849300, 1849899, 1850014, 1850022, 1850023, 1850024, 1850145, 1850471,

Appendix

Here is a scatter plot showing $\log(\text{vruntime})$ vs. sum_exec_runtime ; there is a weak linear correlation.



Figure 9: log vruntime vs. sum exec runtime

Here is a representation of the changes from c3.c to my animation code powered by AI and

the prompt “Show the differences between these two code segments using the formatting of a more reader friendly git diff”

```
Line 3:  
+ #include <string.h>  
  
Lines 8-12:  
+ unsigned p;  
+ char final_strings[4][2048];  
+ for (int i = 0; i < 10; i++) {  
+     memset(final_strings, 0, sizeof(final_strings));  
+     system("top -b -n 1 | tail -n +8 | sort > out.top");  
  
Lines 14-15:  
- for (int i = 0; i < 7; i++) fgets(buffer, sizeof(buffer), fp);  
  
Lines 17-19:  
- sscanf(buffer, " %s %s\n",  
-             pid, user, pr, ni, virt, res, shr, s, cpu, mem, time, command);  
- printf("PID: %s\tTIME: %s\n", pid, time);  
  
Lines 20-27:  
+ sscanf(buffer, " %s %s %s %s %s %s %s %s %s %d\n", pid, user, pr,  
+         ni, virt, res, shr, s, cpu, mem, time, command, &p);  
+ char pid_buf[32];  
+ sprintf(pid_buf, "%s, ", pid);  
+ strcat(final_strings[p], pid_buf);  
+ }  
+ for (int p = 0; p < 4; p++) {  
+     printf("Core %d: %s\n\n", p, final_strings[p]);  
  
Lines 30-32:  
+ system("sleep 1");  
+ system("clear");  
+ }
```